


Storage

Ryan Tolboom

New Jersey Institute of Technology



Mounting in *nix

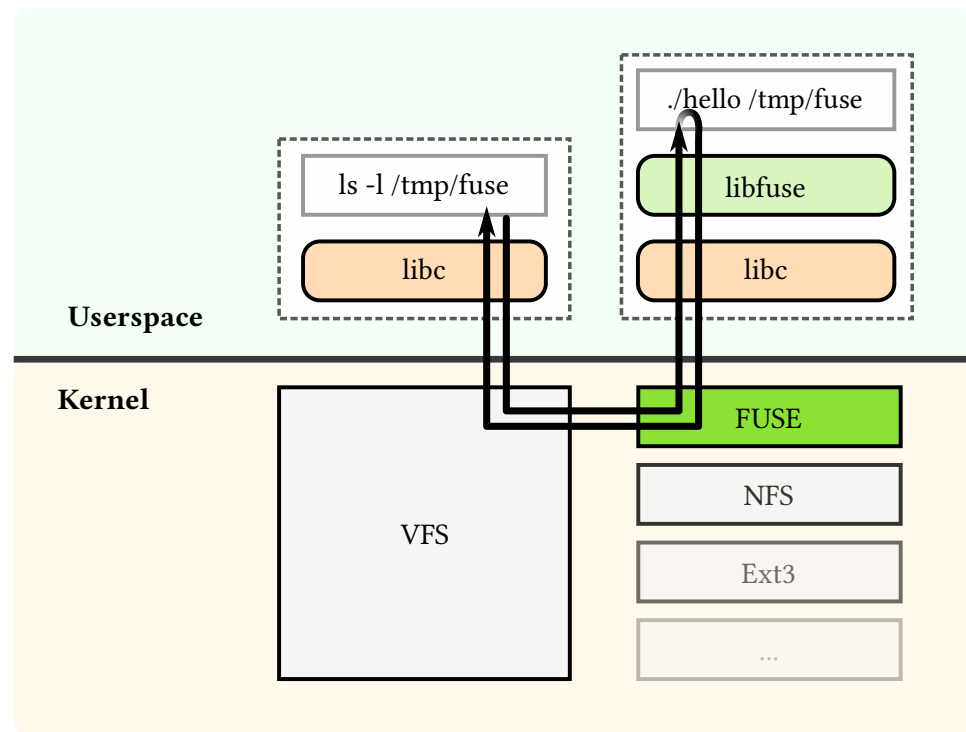


[at least mac is *nix](#) is used under fair use

- All sorts of things can be mounted at points in the file system
- Checkout the output of the mount command
- df shows disk usage for currently mounted things
- /etc/fstab shows you what things are typically mounted at boot
- Things can be mounted on top of one another!
- Block devices are typically mounted in /mnt or /media
- Docker will bind mount directories or mount volumes INTO your containers


Filesystems

- Kernel modules can be loaded to support file systems
- FUSE - file systems can also be setup in userspace
- What is currently supported? (cat /proc/filesystems)
- We often think of filesystems on block devices
- lsblk will show you your block devices
- fdisk (or cfdisk) can be used to see the partition scheme of a disk
 - /dev/sda is the disk
 - /dev/sda1 is the partition
 - /dev/sdb is the NEXT disk
 - YES YOU CAN DESTROY EVERYTHING



["FUSE structure"](#) by [User:Sven](#) is licensed under [CC BY-SA 3.0](#)

Logical Volume Manager




- Modern Linux systems utilize the Logical Volume Manager
- More flexible than traditional partitioning
- Allows for snapshots, encryption, and resizing

"Lvm" by [Emanuel Duss](#) is licensed under [CC BY-SA 3.0](#)

Filesystems in Docker

- Docker uses an overlay fs (implementation differs depending on underlying OS)
- The filesystem a container sees is the thin, top-most layer
- Copy-on-Write (COW) means layers can be shared and copies are only made when needed



Common Filesystems by OS

DOS

- FAT12
- FAT16
- FAT32
- VFAT (widely supported)

Windows

- NTFS
- ReFS

Linux

- ext2
- ReiserFS
- ext3
- ext4
- btrfs
- zfs

Key Features

- scalability
- cryptography
- resilience (CoW, atomic transactions)
- replication

Network File Systems

- SMB
- NFS
- AFS (you should know this one)
- Usually a daemon but can be implemented in the kernel or via FUSE


Common Problems

- keeping things in sync
- permissions
- minimizing bandwidth usage
- caching



[“half filled server racks”](#) by [Alexis Lê-Quôc](#) is licensed under [CC BY-SA 2.0](#)

RAID



[“RAID_5”](#) by [Cburnett](#) is licensed under [CC BY-SA 3.0](#)

Redundant Array of Inexpensive/ Independent Disks

- RAID 0: Striping for speed, make sure your controller can handle the bps
- RAID 1: Redundancy for high availability. Make sure your controller can handle hot-swapping
- RAID 5: Striping AND redundancy. Minimum three disks. A SINGLE lost disk can be recovered from but it takes a while to rebuild the array.
- RAID 1 + 0: Combination of RAID 0 and 1 for speed AND redundancy. More resilient than RAID 5 but less efficient usage of space.

Questions for Discussion

- What's the difference between a file system and a database?
- Are certain file systems better for databases?
- Why do you need to know about networked file systems when using containers?
- Why does Docker use different overlay solutions for different underlying filesystems?



[Question Mark on Chalk Board](#) is in the public domain under [CC0](#)

Resources

- [Understanding file systems](#)
- [Resilient File System \(ReFS\) overview](#)
- [FreeBSD Handbook: Chapter 19. The Z File System \(ZFS\)](#)
- [Wikipedia List of Network File Systems](#)
- [RAID Calculator](#)